

CS 239 Exercise 2 Report

Gerry Agluba Jr.

CS 239 WZZQ- Parallel Computing
Date of Submission: March 26, 2025

1. OBJECTIVES

This report highlights demonstration of (a) improving performance of parallel programs by exploiting different memory strategies namely: shared memory access, access times, patterns, and latency tolerance (b) ability to design reasonable parallel programs based on hardware memory limitations. This exercise revolves around implementing a matrix multiplication program with two matrices of arbitrary shapes inputs leveraging different types of memory access paradigms.

2. METHODOLOGY

The following methods lies on the assumption that a heterogeneous computer system is already setup. In particular, a working CUDA [Cor25a] environment with compatible C++ compiler should be set up prior to design, programming and tests.

2.1 Environment GPU Device Properties

Initially, different device properties were queried to plan different experimental setup that can be made as well as determining the limitations of current environment. A utility function to print device properties were written with the following values observed in table 1

2.2 Kernel Design and Implementation ¹

As stated, two programs of matrix multiplication will be implemented. Kernel function named **matmul_rec_glob** is implemented which calculates the product of two matrices $\mathbf{A} \in \mathbb{R}^{n \times k}$ and $\mathbf{B} \in \mathbb{R}^{k \times m}$ producing an output $\mathbf{C} \in \mathbb{R}^{n \times m}$ implemented using global memory access. Another kernel **matmul_rec_shard** is implemented using tiling algorithm with shared memory access.

For both kernels function; **thread block** of size **32 x 32** is fixed with **TILE_WIDTH** also fixed at the same size. Blocks and tile of said dimensions exactly fits the **1024 max threads per SM** and takes advantage of full utilization of **warp size (32)** and maximization of threads per block.

2.3 Experimental Setup

2.3.1 Experiments

The following items summarized the experiment setup for this report.

- Each experiment will be replicated 10 times.

¹The whole implementation can be found on github repository

Device Count	1
NVIDIA GeForce GTX 1650 SUPER Properties	
Clock Rate (KHz)	1740000
Memory Clock Rate (KHz)	6001000
Total Global Memory (MB)	4095
Shared Memory / Block (KB)	48
Warp Size	32
Pitch (MB)	2047
Max Threads / Block	1024
Max Dimension Size of Block	1024, 1024, 64
Max Dimension Size of Grid	2147483647, 65535, 65535
Multiprocess Count	20
Max Blocks / Multiprocessor	16
Concurrent Kernels	1
Max Threads / Multiprocessor	1024
Shared Memory (KB) / Multiprocessor	64

Table 1: NVIDIA GeForce GTX 1650 SUPER Properties

- Each experiment input matrices will be randomly generated from a uniform distribution $M[i][j] \sim \mathcal{U}(0, 100)$ with matrix size parameters: **n,k and m** all sampled from normal distribution $\mathcal{N}(\mu = \mathbf{b}, \sigma = 0.1\mathbf{b}, \mathbf{b} \in \{2^4, 2^5, 2^6, 2^7, 2^8, 2^9, 2^{12}, 2^{13}\})$ providing a more robust rectangular matrix shapes.
- For this experiment, necessary metrics primarily **runtime (s)** and **compute throughput (%)** will be acquired and sample average over 10 runs will be computed.
- Profiling and Performance Monitoring Tool **NVIDIA Nsight Compute [Cor25b]** will also be utilized for main and other supporting metrics.

3. RESULTS

3.1 Runtime and GPU Throughput

Figure 1 and 3.1 details the kernel runtime results and gpu throughput of both kernel implementations. Runtime seems

to show a comparable runtime performances in both implementations with the `matmul_rec_glob` kernel function inching over `matmul_rec_shar`. On the other hand, a notable performance difference can be seen in the trends of GPU compute throughput, as the tiled matrix multiplication with shared memory access significantly better than the naive implementation.

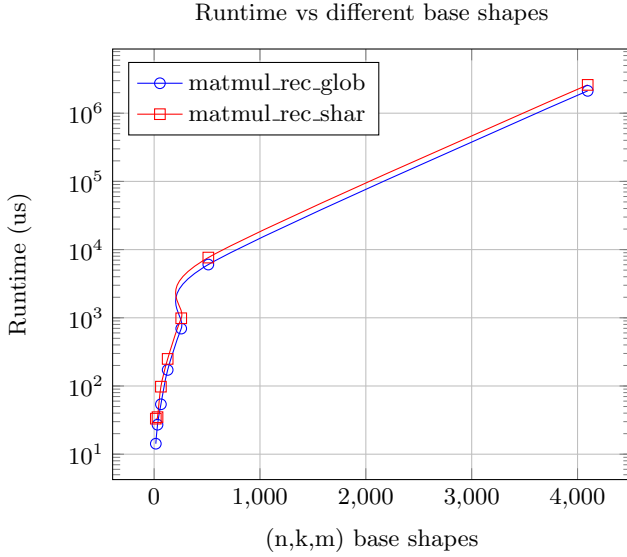


Figure 1: Relative runtime of both kernel implementations: `matmul_rec_glob` and `matmul_rec_shar` over 10 runs acquired from NVIDIA Nsight Compute

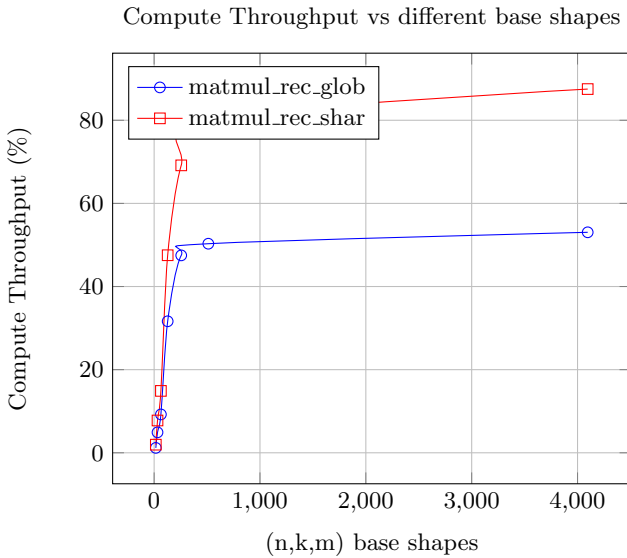


Figure 2: Comparison between the GPU throughputs of `matmul_rec_glob` and `matmul_rec_shar` over 10 runs acquired from NVIDIA Nsight Compute

3.2 Additional NSight Results

An instance of the experiment profiling result is presented below to show empirical observations of both kernel imple-

mentations. `matmul_rec_glob` is labelled in the following results as *baseline* (mostly represented in green color) and *current* is labelled unto `matmul_rec_shar` (mostly represented in blue color). The sample instance is taken from the experiment with the highest matrix size setup.

Figure 3 indicates a +65.41% relative increase on **Compute (SM) Throughput** of `matmul_rec_shar` over `matmul_rec_glob` but worse performance can be seen for **Memory Throughput** with -12.32% and Runtime (+22.25%). Additionally, `matmul_rec_shar` exhibits a **Very High Utilization** of SM ALU with 87% compared `matmul_rec_glob` of 53.1% and have busier SM (+65.37%) as shown in figure 4.




Figure 3: A sample snapshot of the throughput analysis of `matmul_rec_glob` vs `matmul_rec_shar`



Figure 4: A sample snapshot of the compute workload analysis of `matmul_rec_glob` vs `matmul_rec_shar`

On the contrary, `matmul_rec_shar` performs worse with regards to memory through -19.95% relative decrease 5.




Figure 5: A sample snapshot of the memory workload analysis of `matmul_rec_glob` vs `matmul_rec_shar`

Shown in 6, Both kernel functions has achieved high marks on occupancy with almost 100% as well as fully utilizing all warps per SM.

3.3 Computations per Global Memory Access (CGMA)




Figure 6: A sample snapshot of the occupancy of `matmul_rec_glob` vs `matmul_rec_shar`

In the subsection, we calculate both kernels' CGMA.

For `matmul_rec_glob`; each thread in a block requires row(column) of matrix inputs **A** and **B** of size **k** totalling to **2k**. It performs multiplication and addition of single precision (fp32) floating point operations on these memory access aggregating to **2k** operations. Computing CGMA will yield a ratio of $\frac{2k}{2k} = 1$.

On the other hand, kernel `matmul_rec_shared` will load **2** element of associated row(col) of inputs **A** and **B** in collaboration with other threads in a block. Succeeding this is a total operation of $2 * \text{TILE_WIDTH} = 2 * 32$ addition and multiplication operations yielding to a CGMA of $\frac{2 * \text{TILE_WIDTH}}{2} = \text{TILE_WIDTH} = 32$.

Theoretical computations shows **significantly better** CGMA for tiled implementation with shared memory access by a factor of **TILE_WIDTH**.

4. OBSERVATIONS AND NSIGHT (NO PUN INTENDED)

In this report we explore two implementations of matrix multiplication with different memory access strategies. Below are the following insights and observations that were gathered while implementing the program, computing theoretical and analysing empirical results.

4.1 Notes on `matmul_rec_shar` high throughput

Implementations on shared memory enables higher compute efficiency because of data loaded in a shared memory are reused by multiple computations. Computed theoretical CGMA as well as empirical results on ALU utilization manifests this inherent property.

4.2 Arguments on Insignificant Runtime Speed

Contrary to general notion of better performance on shared memory access techniques. Results shows that there are no significant difference between the wall clock runtime of both kernels. These can be attributed to several factors.

Physical Device Shared Memory Capacity can have a significant impact on the runtime of `matmul_rec_shar`; an homage to the ideas presented on **Chapter 5.4 - Memory as a Limiting Factor to Parallelism**[KH12]. Physical shared memory limit of a single SM is **64KB**. Say we have two 4096×4096 single precision float matrix to load. Assuming we use all **20** SM on our current setup. That would take around $\frac{2 \times 4096 \times 4096 \text{ Kb}}{20 \times 64 \text{ Kb}} \approx 102$ cycles to load. If shared memory access is around that **100x** faster or lower, its not suprising that the empirical results we have for runtime speed is not sufficiently significant. This is not yet considering other factors like other memory access latency, thread synchronization etc.

As mentioned, excessive use of `__syncthreads()` can also be one factor though on this report, the influence is not yet measured.

5. RECOMMENDATION

The following items were not yet explored and is pushed as candidate for recommendation that are relevant to this report's objectives.

1. Coalesced Implementation Measures
2. Thread Synchronization
3. Memory Access Latency Observations
4. Other relevant metrics used by NVIDIA Nsight compute
5. Higher variation of matrix sizes
6. Run on environment with **higher shared memory capacity**

6. REFERENCES

References

- [Cor25a] NVIDIA Corporation. *CUDA C++ Programming Guide*. 2025. URL: <https://docs.nvidia.com/cuda/cuda-c-programming-guide/> (visited on 02/28/2025).
- [Cor25b] NVIDIA Corporation. *NVIDIA Nsight Compute*. 2025. URL: <https://developer.nvidia.com/nsight-compute>.
- [KH12] Kirk and Hwu. *Programming Massively Parallel Processors A Hands-on Approach*. Morgan Kaufmann, 2012.

APPENDIX

Listing 1: Kernel Function Implementation

```
--global__ void matmul_rec_glob(float* A, float* B, float* C, int n, int k, int m) {
    // this thread computes C[x][y];
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;

    if (row < n && col < m) {
        float cVal = 0;
        for (int _k = 0; _k < k; _k++) {
            cVal += A[row * k + _k] * B[_k * m + col];
        }
        C[row * m + col] = cVal;
    }
}

--global__ void matmul_rec_shar(float* A, float* B, float* C, int n, int k, int m) {
    __shared__ float ds_A[TILE_WIDTH][TILE_WIDTH];
    __shared__ float ds_B[TILE_WIDTH][TILE_WIDTH];

    int bx = blockIdx.x; int by = blockIdx.y;
    int tx = threadIdx.x; int ty = threadIdx.y;

    int row = by * TILE_WIDTH + ty;
    int col = bx * TILE_WIDTH + tx;

    if (row < n && col < m) {
        float pValue = 0; // partial sum
        for (int q = 0; q < (k + TILE_WIDTH - 1) / TILE_WIDTH; q++) {
            // Load collaboratively a row/col on a submatrix
            ds_A[ty][tx] = A[(row * k) + (q * TILE_WIDTH + tx)];
            ds_B[ty][tx] = B[col + (q * TILE_WIDTH + ty) * k];
            // ds_B[ty][tx] = B[(q * TILE_WIDTH + ty)*m + col]; // more coalesced??

            __syncthreads();
            for (int r = 0; r < TILE_WIDTH; r++) {
                pValue += ds_A[ty][r] * ds_B[r][tx];
                // pValue += ds_A[ty][r] * ds_B[tx][r];
            }
            __syncthreads();
        }
        C[row * m + col] = pValue;
    }
}
```